

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Lasso, Ridge
```

```
In [3]: from sklearn.linear_model import ElasticNet
from sklearn import metrics
```

```
In [4]: df=pd.read_csv("10_USA_Housing.csv")
df
```

Out[4]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenue\nDanielstown, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386
...
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06	USNS Williams\nFPO AP 30153-7653
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06	PSC 9258, Box 8489\nAPO AA 42991- 3352
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06	4215 Tracy Garden Suite 076\nJoshualand, VA 01...
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06	USS Wallace\nFPO AE 73316
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George Ridges Apt. 509\nEast Holly, NV 2...

5000 rows × 7 columns

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                      5000 non-null   float64
1   Avg. Area House Age                   5000 non-null   float64
2   Avg. Area Number of Rooms             5000 non-null   float64
3   Avg. Area Number of Bedrooms          5000 non-null   float64
4   Area Population                       5000 non-null   float64
5   Price                                 5000 non-null   float64
6   Address                               5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

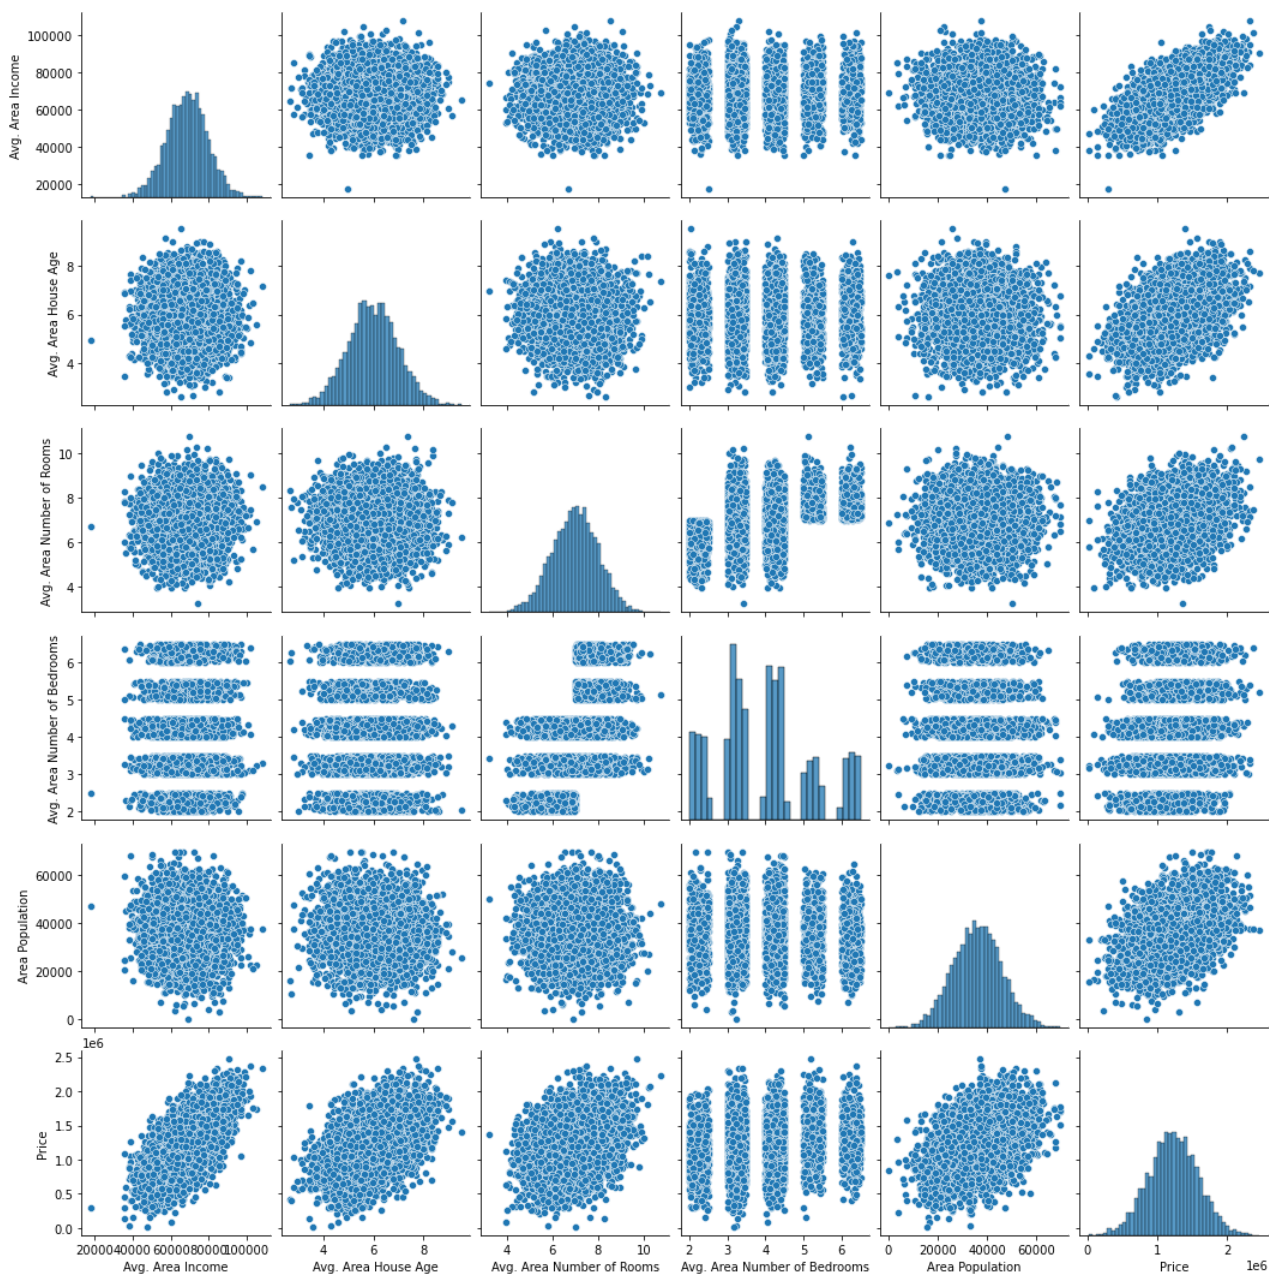
In [6]: `df.describe()`

Out[6]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
min	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
25%	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
50%	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
75%	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
max	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06


```
In [7]: sns.pairplot(df)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x2a536e93070>
```



```
In [8]: sns.displot(df['Price'])
```

```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x2a5382d15e0>
```



```
In [9]: df1=df.drop(['Address'],axis=1)
df1
```

```
Out[9]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05
...
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06

5000 rows × 6 columns

```
In [10]: sns.heatmap(df1.corr())
```

```
Out[10]: <AxesSubplot:>
```



```
In [11]: y=df['Price']
x=df1.drop(['Price'],axis=1)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [12]: model=LinearRegression()
model.fit(x_train,y_train)
model.intercept_
```

```
Out[12]: -2624991.144460169
```


```
In [13]: coeff=pd.DataFrame(model.coef_,x.columns,columns=["Coefficient"])
coeff
```

```
Out[13]:
```

	Coefficient
Avg. Area Income	21.514407
Avg. Area House Age	164315.852135
Avg. Area Number of Rooms	119839.139324
Avg. Area Number of Bedrooms	2503.010663
Area Population	15.272390

```
In [14]: prediction=model.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[14]: <matplotlib.collections.PathCollection at 0x2a538999f40>



```
In [15]: model.score(x_test,y_test)
```

Out[15]: 0.9212486361174832

```
In [16]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
la=Lasso(alpha=10)
la.fit(x_train,y_train)
print(rr.score(x_test,y_test))
la.score(x_test,y_test)
```

0.9212019555165926

Out[16]: 0.9212480013228878

```
In [17]: en=ElasticNet()
en.fit(x_train,y_train)
print(en.coef_)
print(en.intercept_)
print(en.predict(x_test))
print(en.score(x_test,y_test))
```

```
[2.13520337e+01 1.09116497e+05 7.53523245e+04 1.49000056e+04
 1.50889073e+01]
-2016355.419140495
[ 926559.3353008  947524.30558017  979581.46430639 ... 1704221.47260598
 1329984.976759  1272854.93285785]
0.8832584113113195
```

```
In [18]: print("MAE",metrics.mean_absolute_error(y_test,prediction))  
         print("MSE",metrics.mean_squared_error(y_test,prediction))  
         print("RMSE",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

MAE 82276.85863634048

MSE 10469741838.554108

RMSE 102321.75642821085

```
In [19]: df2=pd.read_csv("9_bottle.csv")
df2
```

```
C:\ProgramData\Anaconda3\lib\site-packages\IPython\core\interactiveshell.py:3165:
DtypeWarning: Columns (47,73) have mixed types.Specify dtype option on import or s
et low_memory=False.
    has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
```


Out[19]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...
0	1	1	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0000A-3	0	10.500	33.4400	NaN	25.64900	NaN	...
1	1	2	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0008A-3	8	10.460	33.4400	NaN	25.65600	NaN	...
2	1	3	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0010A-7	10	10.460	33.4370	NaN	25.65400	NaN	...
3	1	4	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0019A-3	19	10.450	33.4200	NaN	25.64300	NaN	...
4	1	5	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0020A-7	20	10.450	33.4210	NaN	25.64300	NaN	...
...
864858	34404	864859	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0000A-7	0	18.744	33.4083	5.805	23.87055	108.74	...
864859	34404	864860	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0002A-3	2	18.744	33.4083	5.805	23.87072	108.74	...
864860	34404	864861	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0005A-3	5	18.692	33.4150	5.796	23.88911	108.46	...
864861	34404	864862	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0010A-3	10	18.161	33.4062	5.816	24.01426	107.74	...

Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...
864862	34404	864863	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0015A-3	15	17.533	33.3880	5.774	24.15297	105.66 ...

864863 rows × 74 columns

In [20]: `df2.isna().sum()`

```
Out[20]: Cst_Cnt          0
Btl_Cnt          0
Sta_ID           0
Depth_ID         0
Depthm           0
          ...
TA1              862779
TA2              864629
pH2              864853
pH1              864779
DIC Quality Comment 864808
Length: 74, dtype: int64
```

In [21]: `df2.columns`

```
Out[21]: Index(['Cst_Cnt', 'Btl_Cnt', 'Sta_ID', 'Depth_ID', 'Depthm', 'T_degC',
'Salnty', 'O2ml_L', 'STheta', 'O2Sat', 'Oxy_μmol/Kg', 'BtlNum',
'RecInd', 'T_prec', 'T_qual', 'S_prec', 'S_qual', 'P_qual', 'O_qual',
'SThtaq', 'O2Satq', 'ChlorA', 'Chlqua', 'Phaeop', 'Phaqua', 'PO4uM',
'PO4q', 'SiO3uM', 'SiO3qu', 'NO2uM', 'NO2q', 'NO3uM', 'NO3q', 'NH3uM',
'NH3q', 'C14As1', 'C14A1p', 'C14A1q', 'C14As2', 'C14A2p', 'C14A2q',
'DarkAs', 'DarkAp', 'DarkAq', 'MeanAs', 'MeanAp', 'MeanAq', 'IncTim',
'LightP', 'R_Depth', 'R_TEMP', 'R_POTEMP', 'R_SALINITY', 'R_SIGMA',
'R_SVA', 'R_DYNHT', 'R_O2', 'R_O2Sat', 'R_SIO3', 'R_PO4', 'R_NO3',
'R_NO2', 'R_NH4', 'R_CHLA', 'R_PHAEO', 'R_PRES', 'R_SAMP', 'DIC1',
'DIC2', 'TA1', 'TA2', 'pH2', 'pH1', 'DIC Quality Comment'],
dtype='object')
```

In [22]: `df3=df2.drop(['DIC2', 'TA1', 'TA2', 'pH2', 'pH1', 'DIC Quality Comment'],axis=1)`

In []:

In []:

```
In [23]: df4=df3.drop(['BtlNum', 'T_qual', 'S_qual', 'O_qual', 'SThtaq', 'NH3uM', 'C14As1',  
                    'C14A1p', 'C14As2', 'C14A2p', 'DarkAs', 'DarkAp', 'MeanAs', 'MeanAp',  
                    'IncTim', 'LightP', 'R_NH4', 'R_SAMP', 'DIC1'],axis=1)
```

```
In [24]: df4.isna().sum()/len(df4)*100
```

```
Out[24]: Cst_Cnt      0.000000
          Btl_Cnt      0.000000
          Sta_ID      0.000000
          Depth_ID    0.000000
          Depthm      0.000000
          T_degC      1.267600
          Salnty      5.475318
          O2ml_L     19.501586
          STheta      6.092179
          O2Sat      23.540029
          Oxy_μmol/Kg 23.540723
          RecInd      0.000000
          T_prec      1.267600
          S_prec      5.475318
          P_qual     22.096910
          O2Satq     74.817168
          ChlorA     73.952869
          Chlqua     26.096272
          Phaeop     73.952984
          Phaqua     26.095809
          P04uM      52.210119
          P04q       47.762131
          SiO3uM     59.058140
          SiO3qu     40.930991
          NO2uM      60.967691
          NO2q       38.779437
          NO3uM      60.987694
          NO3q       38.726365
          NH3q       6.540227
          C14A1q     1.879835
          C14A2q     1.877754
          DarkAq     2.823915
          MeanAq     2.824031
          R_Depth    0.000000
          R_TEMP     1.267600
          R_POTEMP   5.324196
          R_SALINITY 5.475318
          R_SIGMA    6.111488
          R_SVA      6.101660
          R_DYNHT    5.394727
          R_O2       19.501586
          R_O2Sat    22.941784
          R_SI03     59.057215
          R_P04      52.209194
          R_NO3      60.986769
          R_NO2      60.966766
          R_CHLA     73.952406
          R_PHAEO    73.952522
          R_PRES     0.000000
          dtype: float64
```


In []:

In [25]: df4.describe()

Out[25]:

	Cst_Cnt	Btl_Cnt	Depthm	T_degC	Salnty	O2ml_L	
count	864863.000000	864863.000000	864863.000000	853900.000000	817509.000000	696201.000000	812
mean	17138.790958	432432.000000	226.831951	10.799677	33.840350	3.392468	
std	10240.949817	249664.587267	316.050259	4.243825	0.461843	2.073256	
min	1.000000	1.000000	0.000000	1.440000	28.431000	-0.010000	
25%	8269.000000	216216.500000	46.000000	7.680000	33.488000	1.360000	
50%	16848.000000	432432.000000	125.000000	10.060000	33.863000	3.440000	
75%	26557.000000	648647.500000	300.000000	13.880000	34.196900	5.500000	
max	34404.000000	864863.000000	5351.000000	31.140000	37.034000	11.130000	

8 rows × 47 columns



In [26]: df4.shape

Out[26]: (864863, 49)

In [27]: df5=df4.iloc[0:5000,0:]
df5.shape

Out[27]: (5000, 49)

```
In [28]: per=df5.isna().sum()/len(df5)*100  
per1=pd.DataFrame(per,df5.columns,)  
per1
```

Out[28]:

	0
Cst_Cnt	0.00
Btl_Cnt	0.00
Sta_ID	0.00
Depth_ID	0.00
Depthm	0.00
T_degC	0.40
Salnty	3.04
O2ml_L	43.80
STheta	3.34
O2Sat	45.74
Oxy_μmol/Kg	45.74
Reclnd	0.00
T_prec	0.40
S_prec	3.04
P_qual	0.00
O2Satq	52.88
ChlorA	100.00
Chlqua	0.00
Phaeop	100.00
Phaqua	0.00
PO4uM	79.08
PO4q	20.92
SiO3uM	100.00
SiO3qu	0.00
NO2uM	100.00
NO2q	0.00
NO3uM	100.00
NO3q	0.00
NH3q	0.00
C14A1q	0.00
C14A2q	0.00
DarkAq	0.00
MeanAq	0.00
R_Depth	0.00
R_TEMP	0.40
R_POTEMP	4.50

	0
R_SALINITY	3.04
R_SIGMA	5.62
R_SVA	5.62
R_DYNHT	4.28
R_O2	43.80
R_O2Sat	46.20
R_SIO3	100.00
R_PO4	79.08
R_NO3	100.00
R_NO2	100.00
R_CHLA	100.00
R_PHAEO	100.00
R_PRES	0.00

```
In [29]: pr=per1[per1[0]>75].index  
pr
```

```
Out[29]: Index(['ChlorA', 'Phaeop', 'P04uM', 'Si03uM', 'NO2uM', 'NO3uM', 'R_SIO3',  
               'R_PO4', 'R_NO3', 'R_NO2', 'R_CHLA', 'R_PHAEO'],  
              dtype='object')
```



```
In [30]: df6=df5.drop(['ChlorA', 'Phaeop', 'PO4uM', 'SiO3uM', 'NO2uM', 'NO3uM', 'R_SI03',  
                    'R_PO4', 'R_NO3', 'R_NO2', 'R_CHLA', 'R_PHAE0'],axis=1)  
df6.isna().sum()/len(df5)*100
```

```
Out[30]: Cst_Cnt          0.00  
Btl_Cnt          0.00  
Sta_ID           0.00  
Depth_ID         0.00  
Depthm           0.00  
T_degC           0.40  
Salnty           3.04  
O2ml_L           43.80  
STheta           3.34  
O2Sat            45.74  
Oxy_μmol/Kg       45.74  
RecInd           0.00  
T_prec           0.40  
S_prec           3.04  
P_qual           0.00  
O2Satq           52.88  
Chlqua           0.00  
Phaqua           0.00  
P04q             20.92  
SiO3qu           0.00  
NO2q             0.00  
NO3q             0.00  
NH3q             0.00  
C14A1q           0.00  
C14A2q           0.00  
DarkAq           0.00  
MeanAq           0.00  
R_Depth          0.00  
R_TEMP           0.40  
R_POTEMP         4.50  
R_SALINITY       3.04  
R_SIGMA          5.62  
R_SVA            5.62  
R_DYNHT          4.28  
R_O2             43.80  
R_O2Sat          46.20  
R_PRES           0.00  
dtype: float64
```

In [31]:

```
nul_col = df6.columns[df6.isnull().any()].tolist()
nul_col
```

Out[31]:

```
['T_degC',
 'Salnty',
 'O2ml_L',
 'STheta',
 'O2Sat',
 'Oxy_μmol/Kg',
 'T_prec',
 'S_prec',
 'O2Satq',
 'PO4q',
 'R_TEMP',
 'R_POTEMP',
 'R_SALINITY',
 'R_SIGMA',
 'R_SVA',
 'R_DYNHT',
 'R_O2',
 'R_O2Sat']
```

In [32]:

```
#for i in nul_col:
    #df6[i]= df6.fillna(df6[i].mean())
```

In [33]:

```
#df6
```

In [34]:


```
df6.fillna(df6["T_degC"].mean())
df6.fillna(df6["Salnty"].mean())
df6.fillna(df6["O2ml_L"].mean())
df6.fillna(df6["STheta"].mean())
df6.fillna(df6["O2Sat"].median())
df6.fillna(df6["Oxy_μmol/Kg"].mean())
df6.fillna(df6["T_prec"].mean())
df6.fillna(df6["S_prec"].mean())
df6.fillna(df6["O2Satq"].mean())
df6.fillna(df6["PO4q"].mean())
df6.fillna(df6["R_TEMP"].mean())
df6.fillna(df6["R_POTEMP"].mean())
df6.fillna(df6["R_SALINITY"].mean())
df6.fillna(df6["R_SIGMA"].mean())
df6.fillna(df6["R_SVA"].mean())
df6.fillna(df6["R_DYNHT"].mean())
df6.fillna(df6["R_O2"].mean())
df6=df6.fillna(df6["R_O2Sat"].mean())
```

```
In [35]: df6.isna().sum()
```

```
Out[35]: Cst_Cnt      0
        Btl_Cnt      0
        Sta_ID      0
        Depth_ID     0
        Depthm      0
        T_degC      0
        Salnty      0
        O2ml_L      0
        STheta      0
        O2Sat      0
        Oxy_μmol/Kg  0
        RecInd      0
        T_prec      0
        S_prec      0
        P_qual      0
        O2Satq      0
        Chlqua      0
        Phaqua      0
        P04q      0
        SiO3qu      0
        NO2q      0
        NO3q      0
        NH3q      0
        C14A1q      0
        C14A2q      0
        DarkAq      0
        MeanAq      0
        R_Depth     0
        R_TEMP      0
        R_POTEMP    0
        R_SALINITY  0
        R_SIGMA     0
        R_SVA       0
        R_DYNHT     0
        R_O2        0
        R_O2Sat     0
        R_PRES      0
        dtype: int64
```


```
In [36]: df7=df6.iloc[:10,:10]  
sns.pairplot(df7)
```

```
Out[36]: <seaborn.axisgrid.PairGrid at 0x2a538c09e80>
```



```
In [37]: sns.heatmap(df7.corr())
```

```
Out[37]: <AxesSubplot:>
```



```
In [38]: df6=df6.drop(["Sta_ID", "Depth_ID"],axis=1)
```

```
In [39]: x=df6.drop(["R_O2"],axis=1)
y=df6["R_O2"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
model1=LinearRegression()
model1.fit(x_train,y_train)
model1.intercept_
```

```
Out[39]: 8.562039965909207e-13
```


```
In [40]: coeff=pd.DataFrame(model1.coef_,x.columns,columns=["Coefficient"])  
coeff
```

Out[40]:

	Coefficient
Cst_Cnt	-5.994744e-14
Btl_Cnt	1.913400e-15
Depthm	2.003080e-15
T_degC	-2.119780e-15
Salnty	1.326320e-16
O2ml_L	1.000000e+00
STheta	-5.510215e-15
O2Sat	-4.037433e-16
Oxy_μmol/Kg	1.815876e-15
Reclnd	-4.143030e-16
T_prec	1.748088e-15
S_prec	4.233931e-15
P_qual	-2.636780e-16
O2Satq	-1.707158e-15
Chlqua	-1.110223e-16
Phaqua	-8.326673e-17
PO4q	5.375136e-16
SiO3qu	0.000000e+00
NO2q	1.734723e-18
NO3q	7.589415e-19
NH3q	0.000000e+00
C14A1q	0.000000e+00
C14A2q	0.000000e+00
DarkAq	0.000000e+00
MeanAq	0.000000e+00
R_Depth	4.109804e-17
R_TEMP	-4.442587e-15
R_POTEMP	1.172777e-15
R_SALINITY	-9.341113e-16
R_SIGMA	1.052802e-15
R_SVA	-2.598983e-15
R_DYNHT	-8.467467e-16
R_O2Sat	-8.506138e-16
R_PRES	-2.418154e-15

```
In [41]: pred=model1.predict(x_test)
plt.scatter(y_test,pred)
```

Out[41]: <matplotlib.collections.PathCollection at 0x2a53a420280>



```
In [42]: model1.score(x_test,y_test)
```

Out[42]: 1.0

```
In [43]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
la=Lasso(alpha=10)
la.fit(x_train,y_train)
print(rr.score(x_test,y_test))
la.score(x_test,y_test)
```

0.999999992144969

Out[43]: 0.998824916729663

```
In [44]: en=ElasticNet()
en.fit(x_train,y_train)
print(en.coef_)
print(en.intercept_)
print(en.predict(x_test))
print(en.score(x_test,y_test))
```

```
[ 0.00000000e+00 -9.82451263e-05  1.03329903e-07  0.00000000e+00
 -0.00000000e+00  9.86723609e-01 -0.00000000e+00  0.00000000e+00
 -2.00865323e-04  0.00000000e+00 -0.00000000e+00 -0.00000000e+00
  0.00000000e+00 -7.19474451e-03  0.00000000e+00  0.00000000e+00
 -0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00
  0.00000000e+00  0.00000000e+00  0.00000000e+00  0.00000000e+00
  0.00000000e+00  2.56085023e-05  0.00000000e+00 -0.00000000e+00
 -0.00000000e+00 -0.00000000e+00  6.94836798e-04 -0.00000000e+00
  0.00000000e+00  5.43491997e-05]
0.6473596034558149
[ 2.01446189 50.37821802 50.48878652 ...  0.57884409  5.34386109
  5.85369593]
0.9999840098922101
```



```
In [45]: print("MAE",metrics.mean_absolute_error(y_test,prediction))  
         print("MSE",metrics.mean_squared_error(y_test,prediction))  
         print("RMSE",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
MAE 1232001.5361151793  
MSE 1638369493614.9846  
RMSE 1279988.0833878824
```

```
In [46]: import pickle
```

```
In [49]: file="prediction"  
         model=pickle.dump(model,open(file,'wb'))
```

```
In [ ]:
```

```
In [ ]:
```